

assignment-2

October 30, 2024

```
[1]: import tensorflow as tf
from tensorflow.keras import layers, models, applications
from sklearn.metrics import classification_report, confusion_matrix
import matplotlib.pyplot as plt
import numpy as np
import os
```

```
[2]: import tensorflow as tf

# Define paths (use raw strings for Windows paths to avoid escape sequence
# issues)
train_dir = r'C:\Users\Rohit\OneDrive\Desktop\ROHIT\datasets\train'
validation_dir = r'C:\Users\Rohit\OneDrive\Desktop\ROHIT\datasets\validation'

# Set parameters
img_height, img_width = 224, 224 # ResNet and VGG require 224x224 input size
batch_size = 32

# Data generators for loading and augmenting the images
train_datagen = tf.keras.preprocessing.image.ImageDataGenerator(
    rescale=1./255,
    rotation_range=20,
    width_shift_range=0.2,
    height_shift_range=0.2,
    shear_range=0.2,
    zoom_range=0.2,
    horizontal_flip=True,
    fill_mode='nearest'
)

val_datagen = tf.keras.preprocessing.image.ImageDataGenerator(rescale=1./255)

# Create generators
train_generator = train_datagen.flow_from_directory(
    train_dir,
    target_size=(img_height, img_width),
    batch_size=batch_size,
```

```

        class_mode='categorical'
    )

    validation_generator = val_datagen.flow_from_directory(
        validation_dir,
        target_size=(img_height, img_width),
        batch_size=batch_size,
        class_mode='categorical'
    )

    # Print class indices to verify
    print("Class indices:", train_generator.class_indices)

```

Found 40 images belonging to 2 classes.

Found 40 images belonging to 2 classes.

Class indices: {'cats': 0, 'dogs': 1}

```

[3]: # Load ResNet-50 model without the top layers
base_model_resnet = applications.ResNet50(weights='imagenet',
    ↳include_top=False, input_shape=(img_height, img_width, 3))

# Freeze the base model
base_model_resnet.trainable = False

# Add custom layers on top
model_resnet = models.Sequential([
    base_model_resnet,
    layers.GlobalAveragePooling2D(),
    layers.Dense(256, activation='relu'),
    layers.Dense(train_generator.num_classes, activation='softmax') # Output
    ↳layer
])

model_resnet.compile(optimizer='adam', loss='categorical_crossentropy',
    ↳metrics=['accuracy'])

```

```

[4]: # Load VGG16 model without the top layers
base_model_vgg = applications.VGG16(weights='imagenet', include_top=False,
    ↳input_shape=(img_height, img_width, 3))

# Freeze the base model
base_model_vgg.trainable = False

# Add custom layers on top
model_vgg = models.Sequential([
    base_model_vgg,
    layers.Flatten(),

```

```

        layers.Dense(256, activation='relu'),
        layers.Dense(train_generator.num_classes, activation='softmax') # Output
        layer
    ])

model_vgg.compile(optimizer='adam', loss='categorical_crossentropy',
    metrics=['accuracy'])

```

```

[5]: # Train the ResNet model
history_resnet = model_resnet.fit(
    train_generator,
    steps_per_epoch=train_generator.samples // batch_size,
    validation_data=validation_generator,
    validation_steps=validation_generator.samples // batch_size,
    epochs=10
)

# Train the VGG model
history_vgg = model_vgg.fit(
    train_generator,
    steps_per_epoch=train_generator.samples // batch_size,
    validation_data=validation_generator,
    validation_steps=validation_generator.samples // batch_size,
    epochs=10
)

```

C:\Users\Rohit\AppData\Local\Programs\Python\Python311\Lib\site-packages\keras\src\trainers\data_adapters\py_dataset_adapter.py:121: UserWarning: Your `PyDataset` class should call `super().__init__(**kwargs)` in its constructor. `**kwargs` can include `workers`, `use_multiprocessing`, `max_queue_size`. Do not pass these arguments to `fit()`, as they will be ignored.

```
self._warn_if_super_not_called()
```

Epoch 1/10

1/1 20s 20s/step -

accuracy: 0.6250 - loss: 0.6810 - val_accuracy: 0.5312 - val_loss: 1.0357

Epoch 2/10

1/1 3s 3s/step -

accuracy: 0.5312 - loss: 1.0226 - val_accuracy: 0.3750 - val_loss: 0.7498

Epoch 3/10

C:\Users\Rohit\AppData\Local\Programs\Python\Python311\Lib\contextlib.py:158:

UserWarning: Your input ran out of data; interrupting training. Make sure that your dataset or generator can generate at least `steps_per_epoch * epochs` batches. You may need to use the `.repeat()` function when building your dataset.

```
self.gen.throw(typ, value, traceback)
```

```

1/1          1s 502ms/step -
accuracy: 0.0000e+00 - loss: 0.0000e+00
Epoch 4/10
1/1          3s 3s/step -
accuracy: 0.6250 - loss: 0.6530 - val_accuracy: 0.4688 - val_loss: 0.7149
Epoch 5/10
1/1          3s 3s/step -
accuracy: 0.5312 - loss: 0.6974 - val_accuracy: 0.6250 - val_loss: 0.6498
Epoch 6/10
1/1          0s 66ms/step -
accuracy: 0.0000e+00 - loss: 0.0000e+00
Epoch 7/10
1/1          6s 6s/step -
accuracy: 0.5000 - loss: 0.7504 - val_accuracy: 0.5938 - val_loss: 0.6442
Epoch 8/10
1/1          1s 1s/step -
accuracy: 0.5000 - loss: 0.6790 - val_accuracy: 0.8750 - val_loss: 0.5514
Epoch 9/10
1/1          0s 75ms/step -
accuracy: 0.0000e+00 - loss: 0.0000e+00
Epoch 10/10
1/1          3s 3s/step -
accuracy: 0.3750 - loss: 0.7281 - val_accuracy: 0.5312 - val_loss: 0.6669
Epoch 1/10
1/1          22s 22s/step -
accuracy: 0.2500 - loss: 0.8295 - val_accuracy: 0.5000 - val_loss: 7.5465
Epoch 2/10
1/1          3s 3s/step -
accuracy: 0.5000 - loss: 8.7537 - val_accuracy: 0.6250 - val_loss: 0.6115
Epoch 3/10
1/1          0s 139ms/step -
accuracy: 0.0000e+00 - loss: 0.0000e+00
Epoch 4/10
1/1          14s 14s/step -
accuracy: 0.5625 - loss: 1.1000 - val_accuracy: 0.8125 - val_loss: 0.4216
Epoch 5/10
1/1          3s 3s/step -
accuracy: 0.8750 - loss: 0.4787 - val_accuracy: 1.0000 - val_loss: 0.0592
Epoch 6/10
1/1          0s 38ms/step -
accuracy: 0.0000e+00 - loss: 0.0000e+00
Epoch 7/10
1/1          9s 9s/step -
accuracy: 0.6250 - loss: 0.6814 - val_accuracy: 0.9375 - val_loss: 0.1373
Epoch 8/10
1/1          8s 8s/step -
accuracy: 0.8125 - loss: 0.3458 - val_accuracy: 1.0000 - val_loss: 0.0446
Epoch 9/10

```

```

1/1          0s 39ms/step -
accuracy: 0.0000e+00 - loss: 0.0000e+00
Epoch 10/10
1/1          14s 14s/step -
accuracy: 0.9375 - loss: 0.1650 - val_accuracy: 1.0000 - val_loss: 0.0464

```

```

[6]: # Evaluate ResNet model
val_loss_resnet, val_accuracy_resnet = model_resnet.
    evaluate(validation_generator)
print(f'ResNet Validation Loss: {val_loss_resnet}, Validation Accuracy:
    {val_accuracy_resnet}')

# Get predictions and true labels
y_true_resnet = validation_generator.classes
y_pred_resnet = model_resnet.predict(validation_generator)
y_pred_resnet_classes = np.argmax(y_pred_resnet, axis=1)

# Classification report for ResNet
print("ResNet Classification Report:")
print(classification_report(y_true_resnet, y_pred_resnet_classes))

```

```

2/2          3s 605ms/step -
accuracy: 0.5167 - loss: 0.6674
ResNet Validation Loss: 0.6630467176437378, Validation Accuracy:
0.5249999761581421
2/2          9s 4s/step
ResNet Classification Report:

```

	precision	recall	f1-score	support
0	1.00	0.05	0.10	20
1	0.51	1.00	0.68	20
accuracy			0.53	40
macro avg	0.76	0.53	0.39	40
weighted avg	0.76	0.53	0.39	40

```

[7]: # Evaluate VGG model
val_loss_vgg, val_accuracy_vgg = model_vgg.evaluate(validation_generator)
print(f'VGG Validation Loss: {val_loss_vgg}, Validation Accuracy:
    {val_accuracy_vgg}')

# Get predictions and true labels
y_true_vgg = validation_generator.classes
y_pred_vgg = model_vgg.predict(validation_generator)
y_pred_vgg_classes = np.argmax(y_pred_vgg, axis=1)

```

```
# Classification report for VGG
print("VGG Classification Report:")
print(classification_report(y_true_vgg, y_pred_vgg_classes))
```

```
2/2          8s 2s/step -
accuracy: 1.0000 - loss: 0.0494
VGG Validation Loss: 0.05772433429956436, Validation Accuracy: 1.0
2/2          9s 2s/step
VGG Classification Report:
```

	precision	recall	f1-score	support
0	0.40	0.40	0.40	20
1	0.40	0.40	0.40	20
accuracy			0.40	40
macro avg	0.40	0.40	0.40	40
weighted avg	0.40	0.40	0.40	40

```
[8]: # Plot ResNet accuracy and loss
plt.figure(figsize=(12, 5))

# Accuracy
plt.subplot(1, 2, 1)
plt.plot(history_resnet.history['accuracy'], label='ResNet Train')
plt.plot(history_resnet.history['val_accuracy'], label='ResNet Val')
plt.title('ResNet Model Accuracy')
plt.ylabel('Accuracy')
plt.xlabel('Epoch')
plt.legend()

# Loss
plt.subplot(1, 2, 2)
plt.plot(history_resnet.history['loss'], label='ResNet Train')
plt.plot(history_resnet.history['val_loss'], label='ResNet Val')
plt.title('ResNet Model Loss')
plt.ylabel('Loss')
plt.xlabel('Epoch')
plt.legend()

plt.show()

# Plot VGG accuracy and loss
plt.figure(figsize=(12, 5))

# Accuracy
plt.subplot(1, 2, 1)
```

```

plt.plot(history_vgg.history['accuracy'], label='VGG Train')
plt.plot(history_vgg.history['val_accuracy'], label='VGG Val')
plt.title('VGG Model Accuracy')
plt.ylabel('Accuracy')
plt.xlabel('Epoch')
plt.legend()

# Loss
plt.subplot(1, 2, 2)
plt.plot(history_vgg.history['loss'], label='VGG Train')
plt.plot(history_vgg.history['val_loss'], label='VGG Val')
plt.title('VGG Model Loss')
plt.ylabel('Loss')
plt.xlabel('Epoch')
plt.legend()

plt.show()

```



